

# Melvin Full-System Design

---

Version: draft 0.1

Author: ChatGPT (with Jake)

Purpose: End-to-end specification of how Melvin's binary, graph-based, always-on system should work.

---

## 1. High-Level Goals

---

Melvin is not an LLM with legs. He is a self-modifying, always-on, binary-level graph system whose job is to:

- Survive and perform in unpredictable environments with limited data.
- Learn from very few examples (dozens, not millions) by extracting patterns and reusing them aggressively.
- Integrate all modalities—text, sensors, and motor control—inside one unified substrate.
- Rewrite his own internal “code” (graph topology + patterns), not just tune a fixed set of weights.
- Run continuously, including background “sleep” processes (replay, compression, structural updates) while awake.

Core design principle:

One binary substrate. Everything else is patterns over that substrate.

---

## 2. Substrate: Single Binary Graph

---

The entire mind is one graph, persisted in memory (and on disk via mmap). No separate systems for text, sensors, or motor control.

Core elements:

- Node table: a single array `nodes[]`
- Edge table: a single array `edges[]`
- Pattern tables: `patterns[]` and `pattern_slots[]` for compact pattern templates

### 2.1 Node Types

We use three logical node types, but they all live in one shared `nodes[]` array.

- DATA node:
  - Holds a single byte (0–255).
  - Represents raw binary chunks from any channel (text, sensor, motor, etc.).
  - Has a small vector embedding `h_data` for semantics.
- BLANK node:
  - Represents a variable slot (e.g., A, B, OBJ, TOOL).
  - Byte field encodes the blank ID (A/B/C...).
  - Has a vector embedding `h_blank` that learns what typically fills this slot.
- PATTERN node:
  - Represents a reference to a reusable pattern template.
  - Does not directly hold structure; instead points to a Pattern entry in `patterns[]`.

- Has a vector embedding `h_pattern` encoding where and when this pattern is useful.

All nodes also carry dynamic state such as activation, decay, and timestamps.

## 2.2 Edge Structure

Edges connect any node to any node. Structurally, all edges are the same type; semantics are encoded by flags.

- Edge fields (conceptual):
  - `src`: `NodeId` (index into `nodes[]`)
  - `dst`: `NodeId` (index into `nodes[]`)
  - `w`: weight (importance / influence)
  - `flags`: bitfield describing semantics:
    - `SEQ`: time/sequence relation
    - `CHAN`: channel tagging (text/motor/sensor)
    - `BIND`: variable binding edge
    - `SEM`: higher-level semantic relation
    - `REPLAY`: links used for episodic replay

There is exactly one `edges[]` array; all connectivity is expressed here.

## 2.3 Patterns and Pattern Slots

Patterns represent reusable templates extracted from data. They are not separate graphs; instead, they are compact descriptors stored in:

- `patterns[]`: each Pattern has:
  - `id`
  - `slot_offset`, `slot_count` (range in `pattern_slots[]`)
  - optional pattern-internal edges (for structural patterns)
  - strength (how often it helped / reliability)
- `pattern_slots[]`: each PatternSlot has:
  - `kind` (`LITERAL_DATA` or `BLANK`)
  - `byte` (literal value or blank ID)
  - `role/position` (e.g., sequence index, graph role)

Each `PATTERN` node in `nodes[]` has an `aux` field that stores a `PatternId`. Matching and binding are done by reading from these tables.

All of this lives in the same mmapped “brain” file; the C code only provides generic update rules.

## 3. Channels and Unified Encoding

---

All input/output streams are bytes. To distinguish modalities, we introduce channel nodes, not separate memories.

- Channel nodes (regular nodes with special meaning):
  - `CH_TEXT`
  - `CH_MOTOR`
  - `CH_SENSOR`
  - `CH_VISION`
  - `CH_REWARD`
  - (can be extended)

When a byte arrives on a particular channel:

- We ensure a DATA node exists for that byte value.
- We create sequence edges linking bytes over time (SEQ edges).
- We create CHAN edges from the channel node to each relevant DATA node.

Example:

A CAN motor command frame [ID, DLC, B0...B7] is represented as a sequence of DATA nodes, each connected to CH\_MOTOR via edges with flags=CHAN.

The same byte (e.g., 0xA3) may appear in text and motor, but the channel edges and surrounding patterns differentiate its roles.

---

#### 4. Node State and Activation Dynamics

---

Each node  $i$  has:

- byte: for DATA/BLANK nodes, the binary payload or blank ID.
- kind: DATA, BLANK, or PATTERN.
- $a_i$ : scalar activation (current energy / firing level).
- $h_i$ : small vector embedding (semantic state, e.g., 16–64 dimensions).
- other fields: in-degree, out-degree, first\_out\_edge, timestamps, etc.

Core update idea (per tick):

1. Message aggregation:

For each node  $i$ , gather messages from incoming edges  $j \rightarrow i$ :

$$m_i = \sum_j (w_{ji} * a_j)$$

Optionally, also aggregate vector messages for  $h_i$ :

$$H_i = \sum_j (w_{ji} * h_j)$$

2. Nonlinear update:

$$a_i(t+1) = (1 - \alpha) * a_i(t) + \alpha * f(m_i + bias_i - decay)$$

where  $f$  is a simple nonlinearity (e.g., clipped linear or tanh).

3. Embedding update (slow drift):

$$h_i(t+1) = (1 - \beta) * h_i(t) + \beta * g(H_i)$$

where  $g$  could be normalization and projection back into a bounded range.

4. Normalization / clipping:

Keep activations and embeddings within safe ranges to avoid blowups.

This gives a continuous-time-like system where activations propagate through the graph at each tick.

---

#### 5. Core Loop: Always-On Operation

---

The brain never stops; only the distribution of activity changes.

At a high level, each global tick does:

1. Input ingestion
2. Activation propagation

3. Pattern matching and binding
4. Prediction and output gating
5. Learning and credit assignment
6. Optional replay / structural updates (sleep work)

## 5.1 Input Ingestion

- Collect new bytes from active channels (text, sensors, motor feedback).
- For each byte:
  - Ensure a DATA node exists for its value.
  - Activate that node (set a\_i to a burst value).
  - Connect it via SEQ edges to the previous byte in that channel's stream.
  - Add CHAN edges from the respective channel node.

## 5.2 Propagation

- For all nodes, update activation and embeddings based on inputs from neighbors.
- This spreads information across the graph and enables downstream patterns to see context.

## 5.3 Pattern Matching and Binding

- For each PATTERN node (or each Pattern in patterns[]):
- Compute a match score against relevant sub-regions of the graph.
- Use attention-like mechanisms:
  - patterns act as queries ( $q_P$ ),
  - local context nodes provide keys ( $k_r$ ),
  - compute similarity  $s(P,r) \approx q_P \cdot k_r$ ,
  - turn scores into soft weights  $\alpha_{P,r}$  via softmax or similar.
- Patterns become active relative to how well they explain the current context.
- Binding:
  - For BLANK slots within patterns, find node sets that satisfy the pattern's constraints.
  - Create or update BIND edges between BLANK nodes and candidate DATA nodes.

## 5.4 Prediction and Output Gating

Patterns don't just match; they also propose:

- Next expected inputs (prediction patterns).
- Appropriate motor commands for current context (control patterns).
- When to produce external output (output gating patterns).

Output gating is itself a pattern, e.g.:

- When there is:
    - An active question context,
    - A filled answer slot with high confidence,
    - No conflicting patterns,
- then strongly activate OUT\_TEXT(byte) nodes.

The host code sees high activation on OUT\_TEXT nodes and pushes those bytes to stdout or to motor buses.

## 5.5 Learning and Credit Assignment

Each tick carries prediction and reward signals.

- Prediction:
  - For each stream, the graph predicts upcoming bytes or pattern activations.
  - When reality arrives, compute prediction error.
- Reward:
  - External reward signals (success/failure, task score) are injected as bytes to CH\_REWARD and/or as a scalar into the learning routine.

Local learning rule (conceptual):

- Each edge and pattern keeps an eligibility trace (how much it contributed recently).
- Weight update:  
 $\Delta w_{ij} \propto r * \text{eligibility}_{ij} + \lambda * (\text{prediction\_error\_component})$
- Pattern strength:
  - Increase when pattern improves prediction or control,
  - Decrease when pattern systematically mispredicts or harms reward.

## 5.6 Replay and Structural Updates (Sleep Work)

When there is slack (or continuously in a background budget), Melvin:

- Replays important episodes:
  - Sequences with high surprise, high reward, or high error.
  - Uses REPLAY edges or stored indices to re-activate those traces.
- Runs more expensive pattern induction and consolidation on replayed episodes:
  - Extracts new patterns,
  - Merges similar ones,
  - Prunes unused or harmful edges and patterns.

This is how “sleep while awake” works: front-end processing and background consolidation share the same loop but may be biased to different tasks at different times.

## 6. Ingestion, Episodes, and Channels

---

### 6.1 Chunking and Episodes

While the graph is continuous, learning benefits from episodic structure.

An episode can be defined as:

- A bounded time window where:
  - Context appears (input text/sensory),
  - Melvin produces some actions (motor commands / outputs),
  - The environment returns feedback (sensors, reward).

In practice:

- Episode boundaries can be marked by special channel bytes (e.g., newline, EOS markers, task-specific flags).
- We can track an episode ID per sequence of SEQ edges for credit assignment and replay.

### 6.2 Channels as Context

Channels tell Melvin where bytes came from without separate memory systems.

Examples:

- Text question:
  - Bytes from microphone → processed into text → CH\_TEXT.
- Visual encoding:
  - Compressed or symbolic visual tokens → CH\_VISION.
- Motor commands:
  - CAN/PWM frames → CH\_MOTOR.
- Feedback:
  - Joint angles, force sensors → CH\_SENSOR, CH\_MOTOR\_FEEDBACK.
- Reward:
  - Task success/failure → CH\_REWARD.

Patterns are free to link any combination of channels. For example:

- A “pick up cup” pattern links CH\_VISION (cup features), CH\_MOTOR (motor frames), and CH\_SENSOR (force profile).
- 

## 7. Pattern Induction from Examples

---

### 7.1 Two-Example Rule for Sequences

Given two sequences (e.g., text, motor frames, or mixed):

1. Align them positionally (simple version: same length; more advanced: dynamic alignment).
2. For each position i:
  - If the nodes are the same (or equivalent):
    - Keep that node as a constant in the pattern.
  - If they differ:
    - Create a BLANK slot at that position.
3. Post-process blank slots:
  - If slot i and j always share fillers across examples, unify them as the same variable (A, B, etc.).
4. Store a Pattern describing:
  - Which positions are constants vs blanks.
  - How these positions are linked by SEQ edges.
  - Statistics about fillers per blank.

Example with arithmetic:

- $1+1=2$
- $2+2=4$

Common positions: + and = → constants.

Differing positions: digits → blanks.

Left/right digits always match → unify to A.

Result differs (2 vs 4), tracked as B.

Pattern: [A] + [A] = [B]

### 7.2 Pattern Induction over Graphs (Objects, Tools, etc.)

For richer sensorimotor episodes, each example is a full graph over time.

Algorithm (conceptual):

1. Identify corresponding subgraphs in example 1 and 2 (e.g., the “grip and lift”

portion).

2. For nodes and edges present in both examples in similar roles:

- Mark as constants (e.g., structure of reach → grip → lift, hinge relations of tools).

3. For nodes and edges that differ but play similar roles (e.g., red vs blue cup, scissors vs pliers):

- Mark as blanks with relational constraints (BLANK OBJ, BLANK TOOL).

4. Store the resulting pattern as a template graph with labeled blanks and edge relations.

This is how Melvin can learn “cup-like grasp” from two cups, or “pincer tool use” from scissors and pliers.

### 7.3 Pattern Library Growth

Over time:

- New patterns are added when:

- They provide significant compression (explain many episodes),
- They correlate strongly with success/reward.

- Old patterns are pruned or merged when:

- Rarely used,
- Outperformed by more general descendants.

The pattern library is the “code” written by data.

---

## 8. Pattern Matching and Application

---

Patterns don’t just recognize—they route and control.

### 8.1 Matching

Given a live context (a subgraph/sequence currently active):

- For each pattern P:

- Compute a match score based on:
  - How many constant slots align,
  - How well blank constraints can be satisfied,
  - Similarity in embeddings  $h_i$  between pattern slots and candidate nodes.

- Use attention-like competition:

- Normalize scores among patterns that cover the same region.
- Allow multiple patterns to be partially active if helpful.

### 8.2 Binding

For blanks:

- Find candidates that satisfy relational constraints (edges and roles).

- Create BIND edges between BLANK nodes and candidate DATA nodes.

- The BLANK embeddings  $h_{\text{blank}}$  update to reflect the distribution of fillers.

### 8.3 Using Patterns to Predict and Control

Once P is active with certain bindings:

- P can:

- Predict missing nodes (fill blanks) → “what comes next” or “what object fits here?”.

- Propose motor commands that historically follow this configuration.
- Gate output (when to speak, when to act).

Examples:

- Arithmetic:
    - Pattern  $[A]+[A]=[B]$  with A bound to '3' → predict B as '6'.
  - Cup grasp:
    - Pattern  $P_{cup\_grasp}$  with OBJ bound to a new cup-like visual cluster → propose a motor trajectory previously associated with successful grasps.
  - Pincer tool pattern:
    - Pattern  $P_{tool\_pincer}$  generalizes from scissors to pliers as long as structural BLANKs can be filled consistently.
- 

## 9. Learning: Weights, Patterns, and Structure

---

Learning happens at three levels:

- 1) Weight learning (fast, local).
- 2) Pattern strength / statistics (medium-term).
- 3) Structural learning (slow: new nodes, edges, patterns; pruning).

### 9.1 Weight Learning

Edges and patterns maintain eligibility traces (e.g., decaying values that track recent involvement).

Upon receiving prediction error or reward  $r$ :

- For each edge  $(i,j)$ :
 
$$\Delta w_{ij} = \eta * r * \text{eligibility}_{ij} + \beta * (\text{prediction\_gradient\_component})$$
- Eligibility  $_{ij}$  decays over time and spikes when the edge is active.

This gives:

- Hebbian-like plasticity modulated by reward/prediction error.
- Local, online updates; no global backprop.

### 9.2 Pattern Statistics

Each pattern  $P$  tracks:

- Usage count.
- Success rate (how often its predictions/control actions were correct or rewarded).
- Compression gain (how many episodes it helps explain vs memory cost).

Pattern strength increases when:

- $P$  matches and leads to good outcome or accurate prediction.

Pattern strength decreases when:

- $P$  repeatedly mispredicts or leads to poor reward.

### 9.3 Structural Learning

Structural changes include:

- Creating new patterns from episodes (induction).
- Merging similar patterns (clustering based on slot structure and statistics).
- Pruning patterns and edges that are:
  - Rarely used,
  - Weak,
  - Redundant relative to stronger patterns.

This keeps the graph compact and focused on useful regularities.

---

## 10. Motor, Text, and Sensor Integration

---

Because everything is bytes in one graph, integrating modalities is primarily about patterns over channels.

### 10.1 Motor Control at Binary Level

Motor output path:

- When PATTERNs plus context strongly activate a sequence of DATA nodes linked to CH\_MOTOR:
  - Host code packages the active sequence into a CAN/PWM frame.
  - Frame is sent to actuators.

Motor feedback path:

- Incoming motor telemetry (e.g., joint positions, status) is converted to bytes and fed to CH\_MOTOR\_FEEDBACK / CH\_SENSOR exactly like other streams.

Teaching motor control:

- Episodes contain:
  - Context (vision, text, internal state),
  - Motor commands (CH\_MOTOR),
  - Resulting sensory feedback,
  - Reward.
- Pattern induction + weight learning discover mappings:
  - Context patterns → motor command patterns → good outcomes.

### 10.2 Text and Symbolic Reasoning

Text is just another channel:

- UTF-8 bytes or tokenized representations feed into CH\_TEXT.
- Patterns capture:
  - Arithmetic relations,
  - Language patterns,
  - Question/answer structures.

Because text, sensors, and motor are in one graph, patterns naturally form bridging structures, e.g.:

- "If user says 'raise arm', and environment is safe → apply motor pattern P\_raise\_arm."

---

## 11. Sleep, Replay, and Compression

---

Melvin does not turn off. When external input is sparse or while idle, he:

- Replays selected episodes using REPLAY edges.
- Runs pattern induction over replayed episodes to discover new patterns.
- Consolidates knowledge:
  - Strengthens useful patterns/edges,
  - Prunes or merges weak/redundant ones.

Replay selection criteria:

- High surprise (large prediction errors).
- High reward (successful behaviors).
- High uncertainty (unstable predictions).

This turns a few real-world experiences into many internal training passes without requiring more environment data.

---

## 12. Data Efficiency Principles

---

To work with 1/10th the data of an LLM-like system, Melvin relies on:

1. Structure-first learning:
  - Extract reusable patterns from very few examples by identifying invariants (constants) and blanks (variables).
2. Aggressive sharing:
  - Same DATA nodes reused across all contexts.
  - Patterns and blanks reused across tasks and modalities.
3. Surprise-driven updates:
  - Focus learning on episodes where predictions fail or outcomes are surprising, not on redundant data.
4. Episodic replay:
  - Reuse each real-world episode many times internally.
5. Online adaptation:
  - Always learning; no strict train/freeze phases.

---

## 13. Minimal Non-Learned Seed

---

The only parts that must be hand-designed in C:

- Data structures:
  - Node, Edge, Pattern, PatternSlot, channel nodes.
- Dynamics:
  - Activation update equations.
  - Basic normalization, decay, and noise rules.
- Learning framework:
  - Eligibility traces and local weight updates.
  - Mechanisms for pattern induction, merging, pruning.
- I/O adapters:
  - Byte ingestion from various channels.
  - Mapping OUT\_TEXT/OUT\_MOTOR activations to actual stdout and motor buses.
- Scheduler:
  - Basic loop and resource splitting between live input handling and background replay/compression.

Everything else—arithmetic, language structure, motor skills, “when to speak”, “how to pick up a cup”, “how to use a pincer tool”—is represented by patterns and graph structure that emerge from data.

---

## 14. Example Scenarios

---

### 14.1 Arithmetic: 3+3=?

- Input bytes: "3+3=?  
" on CH\_TEXT.
- SEQ and CHAN edges link them.
- Pattern [A]+[A]=[B] matches and binds A to '3'.
- Lower-level circuits (built from data) map '3' to value(3), compute value(6), map back to '6' DATA node.
- Output gating pattern sees:
  - Question context,
  - Answer slot filled with '6',
  - High confidence,then activates OUT\_TEXT('6').

### 14.2 Picking Up Any Cup

- Two episodes: picking up a red cup and a blue cup.
- Pattern induction finds common reach-grip-lift structure and marks cup-specific details as blanks (OBJ, GRASP\_REGION).
- P\_cup\_grasp pattern encodes “graspable object with certain affordances”.
- Faced with a new cup-like object, Melvin binds OBJ to that object and reuses P\_cup\_grasp motor sequence, with local adjustment guided by feedback.

### 14.3 Scissors to Pliers

- Two episodes: using scissors and using pliers.
- Pattern induction over their sensorimotor graphs finds a shared pincer structure:
  - Two handles, a pivot, jaws applying force.
- Pattern P\_tool\_pincer encodes this relational structure with blanks for TOOL and TARGET.
- Faced with new pliers or other pincer tool, Melvin can bind TOOL to the new object and reuse the action pattern.

---

## 15. Future Extensions

---

Future improvements can include:

- More sophisticated graph alignment for pattern induction (subgraph isomorphism approximations).
- Multiple timescales of memory (fast plastic edges vs slow structural changes).
- Richer internal modulators (e.g., synthetic neuromodulators influencing plasticity).
- Integration with external systems (LLMs) for distillation: use LLMs as teachers, but keep Melvin grounded and structurally adaptive.
- Advanced visualization tools for the running brain (3D node/edge display, activation overlays, pattern firing timelines).

---

End of Document

---